

# Linux Root Filesystems

**Advanced Embedded Linux  
Development**  
with **Dan Walkes**



University of Colorado **Boulder**

**Learning objectives:**

**Minimal Root Filesystems**

**Filesystem Hierarchy Standard**

**Busybox overview**

# Booting Our QEMU Kernel

- QEMU\_AUDIO\_DRV=none qemu-system-arm -m 256M -nographic -M versatilepb -kernel zImage
  - What is missing?

```
[ 1.886379] Kernel panic - not syncing: VFS: Unable to mount root fs on  
unknown-block(0,0)  
[ 1.895105] ---[ end Kernel panic - not syncing: VFS: Unable to mount root  
fs on unknown-block(0, 0)
```

- Panic is because we need a root filesystem!

# Root Filesystem - initramfs

- Option 1 - initramfs
  - set of files extracted into RAM, typically used to setup your real root filesystem
  - May load modules to support hardware needed to access the root filesystem.
  - We will use it for simplicity with QEMU for Assignment 3.

# Root Filesystem

- Option 2 - block device
  - block device supported by the kernel, specified in a kernel command line with a root= parameter

# Minimal Root Filesystem Contents

- **init**
  - starts the system - usually scripts
- **shell**
  - Runs shell scripts, handles command prompt
- **daemons**
  - background programs that provide services

# Minimal Root Filesystem Contents

- shared libraries
- configuration files (/etc)
- device nodes (/dev)
- /proc and /sys pseudo filesystems
- kernel modules (/lib/modules/[kernel version])

# Filesystem Hierarchy Standard

- Where does the kernel expect files to be located in the root filesystem?
  - It doesn't!
  - Only expects the init program, specified in command line
- What about programs?
  - Expectations Defined in the Filesystem Hierarchy Standard (FHS)



# Filesystem Hierarchy Standard

- /bin - programs for all users, used at boot
- /dev - device nodes and other files
- /etc - system configuration files
- /lib - shared libraries
- /proc, /sys - proc and sysfs filesystem
- /sbin - programs for the system administrator, used at boot

# Filesystem Hierarchy Standard

- /tmp - temporary files - can be deleted on boot
- /usr /usr/bin, /usr/sbin - additional programs libraries, utilities
- /var - files modified at runtime (/var/log) which need to be retained after boot

# How do we create a rootfs?

- Start by creating a folder tree

```
mkdir -p bin dev etc home lib lib64 proc sbin sys tmp usr var
mkdir -p usr/bin usr/lib usr/sbin
mkdir -p var/log
```

```
$ tree
.
├── bin
├── dev
├── etc
├── home
├── lib
├── lib64
├── proc
├── sbin
├── sys
├── tmp
├── usr
│   ├── bin
│   ├── lib
│   └── sbin
├── var
│   └── log
└──
```

16 directories, 0 files

# What about the files?

- Now we've got the filesystem structure, where do we get the files we need to fill it?
- Use Busybox
- Single binary that implements essential Linux programs
- Symbolic links are used to tell which program is being requested

# BusyBox Overview

- Single binary that implements essential Linux programs
- Symbolic links are used to tell which program is being requested

```
$ ls -l bin/cat bin/busybox
-rwxr-xr-x 1 root root 1025472 May 30 20:13 bin/busybox
lrwxrwxrwx 1 root root      7 May 30 20:13 bin/cat -> busybox
```

# Creating Root Filesystem

- Make and install busybox:
  - make distclean
  - make defconfig
  - make ARCH=\${ARCH}
  - CROSS\_COMPILE=\${CROSS\_COMPILE}
  - make CONFIG\_PREFIX=/path/to/rootdir  
ARCH=\${ARCH}  
CROSS\_COMPILE=\${CROSS\_COMPILE} install

# Creating Root Filesystem

- Add needed shared libraries from toolchain sysroot

```
$ aarch64-none-linux-gnu-readelf -a bin/busybox | grep "program interpreter"
[Requesting program interpreter: /lib/ld-linux-aarch64.so.1]
$ aarch64-none-linux-gnu-readelf -a bin/busybox | grep "Shared library"
0x0000000000000001 (NEEDED)           Shared library: [libm.so.6]
0x0000000000000001 (NEEDED)           Shared library: [libresolv.so.2]
0x0000000000000001 (NEEDED)           Shared library: [libc.so.6]
```

- Program interpreter placed in “lib” directory
- Libraries placed in lib64 directory (since arch is 64 bit)

# Devices

- Created with mknod (make node)
- `mknod <name> <type> <major> <minor>`
  - Null device is a known major 1 minor 3
  - Console device is known major 5 minor 1

```
sudo mknod -m 666 dev/null c 1 3  
sudo mknod -m 600 dev/console c 5 1
```



# Make the contents owned by root

```
$ cd rootfs/  
$ sudo chown -R root:root *
```

## Why use root as owner?

- Your user account doesn't exist on the system we are creating
  - Note: you need to be sudo root to run this command

# Using your rootfs with the target

- Ramdisk
  - Disk image loaded in RAM by bootloader
- Disk Image
  - Disk image for use with memory (like sdcard)
- Network File System (NFS)
- We will use a Ramdisk for assignment 3.

# Use CPIO to create initramfs

```
cd "$OUTDIR/rootfs"  
find . | cpio -H newc -ov --owner root:root > ${OUTDIR}/initramfs.cpio
```

```
gzip -f initramfs.cpio
```

# Running in qemu

- Run the start-qemu-terminal script
  - References your kernel image and rootfs

You should see a root shell prompt after a boot-up delay